

Faculty of Engineering of the University of Porto



Project 1

Query Optimization

M.EIC028 - Database Technologies (TBD)

1MEIC02

Professor

Mariana Curado Malta

Student

Diogo Alexandre da Costa Melo Moreira da Fonte - up20204175@edu.fe.up.pt

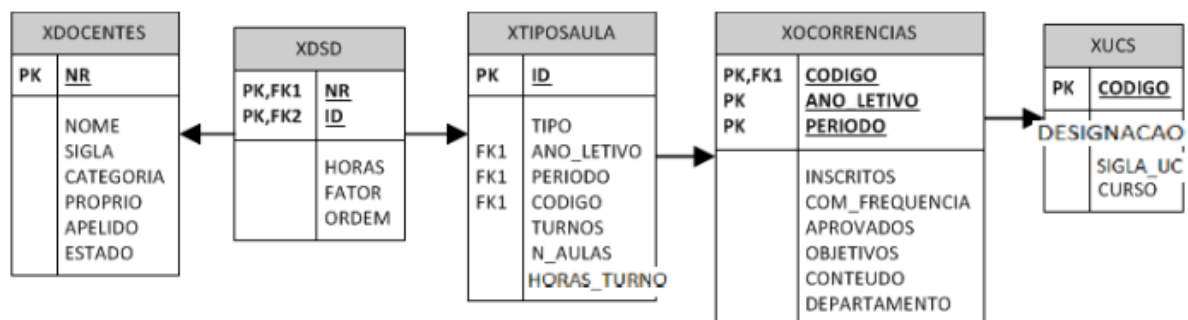
Nuno André Gomes França - up201807530@edu.fe.up.pt

1. Introduction

The purpose of this homework is to analyze several SQL execution plans in a test database. We also need to assess the impact of indexes comparing the statistics of the executions of different strategies for query organization and executed in three different experimentation environments.

The database model we are working with stores data about the distribution of teaching services in a faculty. There are:

- courses (table XUCS);
- occurrences of courses (table XOCORRENCIAS);
- types of classes that belong to occurrences (table XTIPOSAULA) and each occurrence may have one or more class types;
- teaching service distributions (table XDSD) that records, for each class type of occurrence, how many weekly hours are assigned to that professor. If a professor is teaching, in a single class, more than one course at the same time, for example, from different programs, the weight of that course may be less than 1 (attribute factor);
- list of professors (table XDOCENTES).



(figure 1 - Relation Model)

2. Experimentation Environments

We were asked to implement three different experimentation environments, to be able to compare the execution of different queries, analyzing the impact of the query organization and the influence of adding integrity constraints and indexes in our database tables. Below, we explain in more detail each experimentation environment.

2.1. Environment X

In this environment we only created the sql tables, without adding any index or integrity constraint. We created the tables copying the original tables provided.

Here it is the SQL script to create this environment X:

```
CREATE TABLE xdocentes AS SELECT * FROM GTD10.docentes;

CREATE TABLE xdsd AS SELECT * FROM GTD10.dsd;

CREATE TABLE xtiposaula AS SELECT * FROM GTD10.tiposaula;

CREATE TABLE xocorrencias AS SELECT * FROM GTD10.ocorrencias;

CREATE TABLE xucs AS SELECT * FROM GTD10.ucs;
```

2.2. Environment Y

In the environment Y, we added some integrity constraints, primary keys and foreign keys. The constraints are also represented in the relational model above, in the figure 1. We implemented some basic primary and foreign keys, but also some compound keys.

The added primary keys, either compound or not, can uniquely identify each record/row in each table, and the foreign keys are used to represent the relations between the tables, mentioned in the introduction section.

Here it is the SQL script to create this environment Y:

```
CREATE TABLE ydocentes AS SELECT * FROM GTD10.docentes;
ALTER TABLE ydocentes ADD PRIMARY KEY ("NR");

CREATE TABLE yucs AS SELECT * FROM GTD10.ucls;
ALTER TABLE yucs ADD PRIMARY KEY ("CODIGO");

CREATE TABLE yocorrencias AS SELECT * FROM GTD10.ocorrencias;
ALTER TABLE yocorrencias ADD CONSTRAINT YOCORRENCIAS_PK PRIMARY
KEY("CODIGO","ANO_LETIVO", "PERIODO");
ALTER TABLE yocorrencias ADD FOREIGN KEY ("CODIGO") REFERENCES
yucs("CODIGO");

CREATE TABLE ytiposaula AS SELECT * FROM GTD10.tiposaula;
ALTER TABLE ytiposaula ADD PRIMARY KEY ("ID");
ALTER TABLE ytiposaula ADD FOREIGN KEY ("CODIGO", "ANO_LETIVO", "PERIODO")
REFERENCES yocorrencias("CODIGO", "ANO_LETIVO", "PERIODO");

CREATE TABLE ydsd AS SELECT * FROM GTD10.dsd;
ALTER TABLE ydsd ADD CONSTRAINT YDSD_PK PRIMARY KEY("NR","ID");
ALTER TABLE ydsd ADD FOREIGN KEY ("NR") REFERENCES ydocentes("NR");
ALTER TABLE ydsd ADD FOREIGN KEY ("ID") REFERENCES ytiposaula("ID");
```

2.3. Environment Z

Finally, in our environment Z, in addition to adding the integrity constraints mentioned above, in the environment Y description, we added some indexes. Our choice of indexes was based on the analysis of the questions presented in the assignment and we created indexes in the fields that would be most useful to make our queries more efficient. Only one exception, in the table ZTIPOSAULA we tested two types of indexes in the fields "TIPO" and "ANO_LETIVO", as requested in question number five of the assignment.

Here it is the SQL script to create this environment Z:

```
CREATE TABLE zdocentes AS SELECT * FROM GTD10.docentes;
ALTER TABLE zdocentes ADD PRIMARY KEY ("NR");
CREATE INDEX IDX_NOME_DOCENTE ON zdocentes("NOME");

CREATE TABLE zucs AS SELECT * FROM GTD10.ucs;
ALTER TABLE zucs ADD PRIMARY KEY ("CODIGO");
CREATE INDEX IDX_DESIGNACAO_UC ON zucs("DESIGNACAO");
CREATE INDEX IDX_CURSO_UC ON zucs("CURSO");

CREATE TABLE zocorrencias AS SELECT * FROM GTD10.ocorrencias;
ALTER TABLE zocorrencias ADD CONSTRAINT ZOCORRENCIAS_PK PRIMARY
KEY("CODIGO","ANO_LETIVO", "PERIODO");
ALTER TABLE zocorrencias ADD FOREIGN KEY ("CODIGO") REFERENCES
zucs("CODIGO");
CREATE INDEX IDX_DEPARTAMENTO_OCORRENCIA ON zocorrencias("DEPARTAMENTO");

CREATE TABLE ztiposaula AS SELECT * FROM GTD10.tiposaula;
ALTER TABLE ztiposaula ADD PRIMARY KEY ("ID");
ALTER TABLE ztiposaula ADD FOREIGN KEY ("CODIGO", "ANO_LETIVO", "PERIODO")
REFERENCES zocorrencias("CODIGO", "ANO_LETIVO", "PERIODO");
CREATE INDEX BTREE_IDX_TIPO_AND_ANO_LETIVO_TIPOSAULA ON ztiposaula("TIPO",
"ANO_LETIVO"); --BTREE
--CREATE BITMAP INDEX BITMAP_IDX_TIPO_AND_ANO_LETIVO_TIPOSAULA ON
ztiposaula("TIPO", "ANO_LETIVO"); --BITMAP

CREATE TABLE zdspd AS SELECT * FROM GTD10.dspd;
ALTER TABLE zdspd ADD CONSTRAINT ZDSPD_PK PRIMARY KEY("NR","ID");
ALTER TABLE zdspd ADD FOREIGN KEY ("NR") REFERENCES zdocentes("NR");
ALTER TABLE zdspd ADD FOREIGN KEY ("ID") REFERENCES ztiposaula("ID");
```

3. Questions Analysis

In this section, we are going to approach each question of the assignment, analyzing our queries, the retrieved results and the comparison between the executions in the several environments.

3.1. Question 1

Description: Selection and Join - Show the codigo, designacao, ano_letivo, inscritos, tipo, and turnos for the course 'Bases de Dados' of the program 275.

SQL Query:

```
SELECT ucs.codigo, ucs.designacao, ocorrencias.ano_letivo,
ocorrencias.inscritos, tiposaula.tipo, tiposaula.turnos
FROM tiposaula
LEFT JOIN ocorrencias ON tiposaula.codigo = ocorrencias.codigo AND
tiposaula.ano_letivo = ocorrencias.ano_letivo AND tiposaula.periodo =
ocorrencias.periodo
LEFT JOIN ucs ON ocorrencias.codigo = ucs.codigo
WHERE ucs.curso = 275 AND ucs.designacao = 'Bases de Dados'
```

Retrieved Results:

	CODIGO	DESIGNACAO	ANO_LETIVO	INSCRITOS	TIPO	TURNOS
1	EIC3106	Bases de Dados	2003/2004	92	T	1
2	EIC3106	Bases de Dados	2003/2004	92	TP	4
3	EIC3106	Bases de Dados	2004/2005	114	T	1
4	EIC3106	Bases de Dados	2004/2005	114	TP	4
5	EIC3111	Bases de Dados	2005/2006	(null)	T	1
6	EIC3111	Bases de Dados	2005/2006	(null)	TP	6

(figure 2 - Retrieved Results for Question 1)

Execution Data:

Environment X:

Query Execution Time: 0,032 seconds

OPERATION	OBJECT_NAME	OPTIONS	CARDINALITY	COST	
SELECT STATEMENT				1	642
HASH JOIN				1	642
Access Predicates					
AND					
XOCORRENCIAS.CODIGO=XUCS.CODIGO					
XTIPOSAULA.CODIGO=XOCORRENCIAS.CODIGO					
XTIPOSAULA.ANO_LETIVO=XOCORRENCIAS.ANO_LETIVO					
XTIPOSAULA.PERIODO=XOCORRENCIAS.PERIODO					
MERGE JOIN		CARTESIAN		302	49
TABLE ACCESS	XUCS	FULL		1	13
Filter Predicates					
AND					
XUCS.DESIGNACAO='Bases de Dados'					
XUCS.CURSO=275					
BUFFER		SORT		21019	36
TABLE ACCESS	XTIPOSAULA	FULL		21019	36
TABLE ACCESS	XOCORRENCIAS	FULL		21747	593

(figure 3 - Explain Plan of Environment X for Question 1)

Environment Y:

Query Execution Time: 0,031 seconds

OPERATION	OBJECT_NAME	OPTIONS	CARDINALITY	COST	
SELECT STATEMENT				1	55
HASH JOIN				1	55
Access Predicates					
AND					
YTIPOSAULA.CODIGO=YOCORRENCIAS.CODIGO					
YTIPOSAULA.ANO_LETIVO=YOCORRENCIAS.ANO_LETIVO					
YTIPOSAULA.PERIODO=YOCORRENCIAS.PERIODO					
NESTED LOOPS				1	19
NESTED LOOPS				5	19
TABLE ACCESS	YUCS	FULL		1	13
Filter Predicates					
AND					
YUCS.DESIGNACAO='Bases de Dados'					
YUCS.CURSO=275					
INDEX	YOCORRENCIAS_PK	RANGE SCAN		5	1
Access Predicates					
YOCORRENCIAS.CODIGO=YUCS.CODIGO					
TABLE ACCESS	YOCORRENCIAS	BY INDEX ROWID		5	6
TABLE ACCESS	YTIPOSAULA	FULL		21019	36

(figure 4 - Explain Plan of Environment Y for Question 1)

Environment Z:

Query Execution Time: 0,041 seconds

OPERATION	OBJECT_NAME	OPTIONS	CARDINALITY	COST	
SELECT STATEMENT				1	45
HASH JOIN				1	45
Access Predicates					
AND					
ZTIPOSAULA.CODIGO=ZOCORRENCIAS.CODIGO					
ZTIPOSAULA.ANO_LETIVO=ZOCORRENCIAS.ANO_LETIVO					
ZTIPOSAULA.PERIODO=ZOCORRENCIAS.PERIODO					
NESTED LOOPS				1	9
NESTED LOOPS				5	9
TABLE ACCESS	ZUCS	BY INDEX ROWID BATCHED		1	3
Filter Predicates					
ZUCS.CURSO=275					
INDEX	IDX_DESIGNACAO_UC	RANGE SCAN		2	1
Access Predicates					
ZUCS.DESIGNACAO='Bases de Dados'					
INDEX	ZOCORRENCIAS_PK	RANGE SCAN		5	1
Access Predicates					
ZOCORRENCIAS.CODIGO=ZUCS.CODIGO					
TABLE ACCESS	ZOCORRENCIAS	BY INDEX ROWID		5	6
TABLE ACCESS	ZTIPOSAULA	FULL		21019	36

(figure 5 - Explain Plan of Environment Z for Question 1)

Comparison:

Analyzing the three environments explain plans, we can conclude that the addition of integrity constraints (primary keys (PK) and foreign keys (FK)) is essential for query performance. We can see that in the first environment (X) it needed a *Merge Join*, and in the second and third Nested Loops were used to link the rows using the proper PK and FK.

The indexes didn't make a lot of difference, but were useful to select and join the wanted record/row, and improved a bit the cost of the query.

About the time execution, it was not very conclusive because the execution times collected were a lot similar. We noticed that it has grown with the addition of integrity constraints and with the indexes, but we think that it is a good trade between query execution time versus query execution cost.

3.2. Question 2

Description: Aggregation - How many class hours of each type did the program 233 planned in 2004/2005?

SQL Query:

```
SELECT tiposaula.tipo AS TIPO, SUM(tiposaula.horas_turno *
tiposaula.turnos) AS HORAS
FROM tiposaula
LEFT JOIN ocorrencias ON tiposaula.codigo = ocorrencias.codigo AND
tiposaula.ano_letivo = ocorrencias.ano_letivo AND tiposaula.periodo =
ocorrencias.periodo
LEFT JOIN ucs ON ocorrencias.codigo = ucs.codigo
WHERE ucs.curso = 233 AND ocorrencias.ano_letivo = '2004/2005'
GROUP BY tiposaula.tipo;
```

Retrieved Results:

	TIPO	HORAS
1	P	581.5
2	TP	697.5
3	T	308

(figure 6 - Retrieved Results for Question 2)

Execution Data:

Environment X:

Query Execution Time: 0,039 seconds

OPERATION	OBJECT_NAME	OPTIONS	CARDINALITY	COST
SELECT STATEMENT				643
HASH		GROUP BY		5
HASH JOIN				552
Access Predicates				642
AND				
XTIPOSADLA.CODIGO=XOCORRENCIAS.CODIGO				
XTIPOSADLA.ANO_LETIVO=XOCORRENCIAS.ANO_LETIVO				
XTIPOSADLA.PERIODO=XOCORRENCIAS.PERIODO				
HASH JOIN				552
Access Predicates				606
XOCORRENCIAS.CODIGO=XUCS.CODIGO				
TABLE ACCESS	XUCS	FULL		504
Filter Predicates				13
XUCS.CURSO=233				
TABLE ACCESS	XOCORRENCIAS	FULL		1055
Filter Predicates				593
XOCORRENCIAS.ANO_LETIVO='2004/2005'				
TABLE ACCESS	XTIPOSADLA	FULL		1671
Filter Predicates				36
XTIPOSADLA.ANO_LETIVO='2004/2005'				

(figure 7 - Explain Plan of Environment X for Question 2)

Environment Y:

Query Execution Time: 0,032 seconds

OPERATION	OBJECT_NAME	OPTIONS	CARDINALITY	COST
SELECT STATEMENT				50
HASH		GROUP BY		5
HASH JOIN				53
Access Predicates				49
YTIPOSADLA.CODIGO=ITEM_1				
VIEW	SYS.VW_GRF_5			47
TABLE ACCESS	YUCS	FULL		47
Filter Predicates				13
YUCS.CURSO=233				
TABLE ACCESS	YTIPOSADLA	FULL		1106
Filter Predicates				36
YTIPOSADLA.ANO_LETIVO='2004/2005'				

(figure 8 - Explain Plan of Environment Y for Question 2)

Environment Z:

Query Execution Time: 0,029 seconds

OPERATION	OBJECT_NAME	OPTIONS	CARDINALITY	COST
SELECT STATEMENT				43
HASH		GROUP BY		5
HASH JOIN				53
Access Predicates				42
ZTIPOSADLA.CODIGO=ITEM_1				
VIEW	SYS.VW_GRF_5			47
TABLE ACCESS	ZUCS	BY INDEX ROWID BATCHED		47
INDEX	IDC_CURSO_UC	RANGE SCAN		47
Access Predicates				1
ZUCS.CURSO=233				
TABLE ACCESS	ZTIPOSADLA	FULL		1106
Filter Predicates				36
ZTIPOSADLA.ANO_LETIVO='2004/2005'				

(figure 9 - Explain Plan of Environment Z for Question 2)

Comparison:

Comparing the three environments explain plans, we noticed that the implementation of integrity constraints (PK and FK) is crucial for query performance. We can see that in the first environment (X) it needed two Hash Joins, and in the second and third only one. With the proper integrity constraints we can see that in the environments Y and Z it wasn't needed to pass through the table occurrences. It was possible to link directly the table ucs through the attribute codigo in the table tiposaula.

The indexes didn't make a lot of difference, but were useful to select and join the wanted record/row, and improved a bit the cost of the query.

In this question, the execution times were not very conclusive because the times collected were a lot similar. Despite that, we noticed that it has decreased with the addition of integrity constraints and with the indexes, but we think that it is a good trade between query execution time versus query execution cost.

3.3. Question 3

Description: Negation - Which courses (show the code) did have occurrences planned but did not get service assigned in the year 2003/2004?

Retrieved Results for both approaches:

CODIGO	CODIGO	CODIGO	CODIGO	CODIGO
1 MEMT131	16 MEMT102	31 MEMT2000	46 MEA412	61 MPFCA104
2 MEEC1053	17 MEMT107	32 MDI1103	47 MPFCA101	62 MPFCA200
3 MDI1205	18 MEM187	33 MDI1105	48 MPFCA205	63 MEAM5000
4 CI027	19 MEM189	34 MDI1107	49 EIC5127	64 MEMT106
5 MEM157	20 EI1107	35 MDI1108	50 MTM115	65 MDI1106
6 MTM108	21 CI017	36 MDI1208	51 EC5200	66 MTM114
7 MEM181	22 CI007	37 MEM179	52 EC5280	67 MVC1211
8 EIC4220	23 CI014	38 EIC3209	53 EMM528	68 MPPAU1114
9 EIC4222	24 CI008	39 MPPAU1113	54 CI016	69 MPPAU2219
10 EQ418	25 CI018	40 MEA215	55 CI020	70 MEM180
11 MEMT100	26 MEAM1310	41 MPPAU2217	56 CI011	71 CI038
12 MEMT1000	27 MPPAU2215	42 MEA414	57 MPFCA100	72 MEA112
13 MPFCA103	28 MEA219	43 MEAM1312	58 MTM110	73 MEA217
14 MPFCA204	29 MPFCA106	44 MEMT135	59 EIC5124	74 MEA320
15 EIC4221	30 EIC4225	45 MTM111	60 EEC5022	75 EC5287
CODIGO	CODIGO	CODIGO	CODIGO	CODIGO
76 EIC5125	91 MEM163	106 MDI1209	121 MEST210	137 MPFCA102
77 MPFCA105	92 CI009	107 CI002	122 EQ308	138 MPFCA203
78 MPFCA201	93 MEA415	108 CI019	123 MEM5000	
79 MPFCA202	94 MEM184	109 EEC2207	124 EEC5272	
80 MPFCA206	95 CI023	110 GEI512	125 MPPAU1112	
81 EIC5126	96 MPPAU2216	111 MEB204	126 MEM158	
82 MPFCA107	97 MEB205	112 MEAM1314	127 MEM182	
83 MEM188	98 EIC5123	113 MTM104	128 MEM183	
84 MFAMF1108	99 EIC4223	114 EIC5129	129 CI003	
85 EQ407	100 EIC5122	115 EQ411	130 CI004	
86 MEM175	101 MEMT105	116 MMCCE1220	131 CI013	
87 MDI1100	102 CI037	117 EIC4224	132 MEA216	
88 MDI1204	103 CI025	118 MEMT110	133 MEA319	
89 MPPAU2220	104 MEM1205	119 MEMT120	134 MPPAU1115	
90 MEM191	105 MDI1207	120 MDI1206	135 MPPAU2218	

(figure 10 - Retrieved Results for Question 3)

a) Use not in.

SQL Query:

```
SELECT DISTINCT ocorrencias.codigo
FROM ocorrencias
WHERE ocorrencias.ano_letivo = '2003/2004'
AND ocorrencias.codigo NOT IN (
    SELECT DISTINCT ocorrencias.codigo
    FROM dsd
    LEFT JOIN tiposaula ON dsd.ID = tiposaula.ID
    LEFT JOIN ocorrencias ON tiposaula.codigo = ocorrencias.codigo AND
    tiposaula.ano_letivo = ocorrencias.ano_letivo AND tiposaula.periodo =
    ocorrencias.periodo
    WHERE ocorrencias.ano_letivo = '2003/2004'
)
```

Execution Data:

Environment X:

Query Execution Time: 0,102 seconds

OPERATION	OBJECT_NAME	OPTIONS	CARDINALITY	COST
SELECT STATEMENT				1249
HASH		UNIQUE	1	1249
HASH JOIN		RIGHT ANTI	10	1248
Access Predicates XOCORRENCIAS.CODIGO=CODIGO				
VIEW	SYS.VW_NSQ_1		1051	656
HASH JOIN		SEMI	1051	656
Access Predicates XDSO.ID=XTIPOSAULA.ID				
HASH JOIN			1051	629
AND XTIPOSAULA.PERIODO=XOCORRENCIAS.PERIODO XTIPOSAULA.ANO_LETIVO=XOCORRENCIAS.ANO_LETIVO XTIPOSAULA.CODIGO=XOCORRENCIAS.CODIGO				
TABLE ACCESS Filter Predicates XOCORRENCIAS.ANO_LETIVO='2003/2004'	XOCORRENCIAS	FULL	1028	593
TABLE ACCESS Filter Predicates XTIPOSAULA	XTIPOSAULA	FULL	1588	36
Filter Predicates XTIPOSAULA.ANO_LETIVO='2003/2004'				
TABLE ACCESS	XDSO	FULL	27385	27
TABLE ACCESS Filter Predicates XOCORRENCIAS.ANO_LETIVO='2003/2004'	XOCORRENCIAS	FULL	1028	593

(figure 11 - Explain Plan of Environment X for Question 3)

Environment Y:

Query Execution Time: 0,087 seconds

OPERATION	OBJECT_NAME	OPTIONS	CARDINALITY	COST
SELECT STATEMENT				86
HASH		UNIQUE	1	86
HASH JOIN		ANTI	10	85
Access Predicates YOCORRENCIAS.CODIGO=CODIGO				
INDEX	YOCORRENCIAS_PK	FAST FULL SCAN		27
Filter Predicates YOCORRENCIAS.ANO_LETIVO='2003/2004'				
VIEW	SYS.VW_NSQ_1		1588	58
HASH JOIN		SEMI	1588	58
Access Predicates YDSO.ID=YTIPOSAULA.ID				
TABLE ACCESS Filter Predicates YTIPOSAULA.ANO_LETIVO='2003/2004'	YTIPOSAULA	FULL	1588	36
Filter Predicates YTIPOSAULA.ANO_LETIVO='2003/2004'				
INDEX	YDSO_PK	FAST FULL SCAN		22

(figure 12 - Explain Plan of Environment Y for Question 3)

Environment Z:

Query Execution Time: 0,084 seconds

OPERATION	OBJECT_NAME	OPTIONS	CARDINALITY	COST
SELECT STATEMENT				86
HASH		UNIQUE	1	86
HASH JOIN		ANTI	10	85
Access Predicates ZOCORRENCIAS.CODIGO=CODIGO				
INDEX	ZOCORRENCIAS_PK	FAST FULL SCAN		27
Filter Predicates ZOCORRENCIAS.ANO_LETIVO='2003/2004'				
VIEW	SYS.VW_NSQ_1		1588	58
HASH JOIN		SEMI	1588	58
Access Predicates ZDSO.ID=ZTIPOSAULA.ID				
TABLE ACCESS Filter Predicates ZTIPOSAULA.ANO_LETIVO='2003/2004'	ZTIPOSAULA	FULL	1588	36
Filter Predicates ZTIPOSAULA.ANO_LETIVO='2003/2004'				
INDEX	ZDSO_PK	FAST FULL SCAN		22

(figure 13 - Explain Plan of Environment Z for Question 3)

Comparison:

Comparing the execution data of the environments, we noticed a great improvement with the addition of integrity constraints (PK and FK). The cost of queries executed in the environments Y and Z had a cost 1400% smaller than the cost of the query in the environment X.

This time, the indexes didn't make a big difference, the cost remained the same as the cost in the environment Y.

The execution times didn't oscillate much, but it followed a decreasing trend.

b) Use external join, and it is null.

SQL Query:

```
SELECT DISTINCT ucs.codigo
FROM ocorrencias
LEFT JOIN (
    SELECT DISTINCT ocorrencias.codigo AS code
    FROM dsd
    LEFT JOIN tiposaula ON dsd.ID = tiposaula.ID
    LEFT JOIN ocorrencias ON tiposaula.codigo = ocorrencias.codigo AND
    tiposaula.ano_letivo = ocorrencias.ano_letivo AND tiposaula.periodo =
    ocorrencias.periodo
    WHERE ocorrencias.ano_letivo = '2003/2004'
) ON ocorrencias.codigo = code
WHERE ocorrencias.ano_letivo = '2003/2004' and code IS NULL
```

Execution Data:

Environment X:

Query Execution Time: 0,098 seconds

OPERATION	OBJECT_NAME	OPTIONS	CARDINALITY	COST
SELECT STATEMENT				1250
HASH		UNIQUE	1	1250
HASH JOIN		RIGHT ANTI	10	1249
Access Predicates				
XOCORRENCIAS.CODIGO=CODE				
VIEW				940
HASH		UNIQUE	940	657
HASH JOIN		SEMI	1051	657
Access Predicates				
XDSD.ID=XTIPOSADULA.ID				
HASH JOIN			1051	629
Access Predicates				
AND				
XTIPOSADULA.CODIGO=XOCORRENCIAS.CODIGO				
XTIPOSADULA.ANO_LETIVO=XOCORRENCIAS.ANO_LETIVO				
XTIPOSADULA.PERIODO=XOCORRENCIAS.PERIODO				
TABLE ACCESS	XOCORRENCIAS	FULL		1028
Filter Predicates				
XOCORRENCIAS.ANO_LETIVO='2003/2004'				
TABLE ACCESS	XTIPOSADULA	FULL		1588
Filter Predicates				
XTIPOSADULA.ANO_LETIVO='2003/2004'				
TABLE ACCESS	XDSD	FULL		27385
Filter Predicates				
XOCORRENCIAS.ANO_LETIVO='2003/2004'				
TABLE ACCESS	XOCORRENCIAS	FULL		1028
Filter Predicates				
XOCORRENCIAS.ANO_LETIVO='2003/2004'				

(figure 15 - Explain Plan of Environment X for Question 3)

Environment Y:

Query Execution Time: 0,083 seconds

OPERATION	OBJECT_NAME	OPTIONS	CARDINALITY	COST
SELECT STATEMENT			1	87
HASH		UNIQUE	1	87
HASH JOIN		ANTI	10	86
Access Predicates				
YOCORRENCIAS.CODIGO=CODE				
INDEX	YOCORRENCIAS_PK	FAST FULL SCAN		27
Filter Predicates				
YOCORRENCIAS.ANO_LETIVO='2003/2004'				
VIEW			1322	59
HASH		UNIQUE	1322	59
HASH JOIN		SEMI	1588	58
Access Predicates				
YDSD.ID=YTIPOSAULA.ID				
TABLE ACCESS	YTIPOSAULA	FULL	1588	36
Filter Predicates				
YTIPOSAULA.ANO_LETIVO='2003/2004'				
INDEX	YDSD_PK	FAST FULL SCAN	27385	22

(figure 16 - Explain Plan of Environment Y for Question 3)

Environment Z:

Query Execution Time: 0,080 seconds

OPERATION	OBJECT_NAME	OPTIONS	CARDINALITY	COST
SELECT STATEMENT			1	87
HASH		UNIQUE	1	87
HASH JOIN		ANTI	10	86
Access Predicates				
ZOCORRENCIAS.CODIGO=CODE				
INDEX	ZOCORRENCIAS_PK	FAST FULL SCAN		27
Filter Predicates				
ZOCORRENCIAS.ANO_LETIVO='2003/2004'				
VIEW			1322	59
HASH		UNIQUE	1322	59
HASH JOIN		SEMI	1588	58
Access Predicates				
ZDSD.ID=ZTIPOSAULA.ID				
TABLE ACCESS	ZTIPOSAULA	FULL	1588	36
Filter Predicates				
ZTIPOSAULA.ANO_LETIVO='2003/2004'				
INDEX	ZDSD_PK	FAST FULL SCAN	27385	22

(figure 17 - Explain Plan of Environment Z for Question 3)

Comparison:

Comparing the execution data of the three environments, we noticed a great improvement with the addition of integrity constraints (PK and FK). The cost of queries executed in the environments Y and Z had a cost 1400% smaller than the cost of the query in the environment X.

This time, the indexes didn't make a big difference, the cost remained the same as the cost in the environment Y.

The execution times oscillated a bit, but it isn't a good parameter to retrieve more conclusions for this comparison.

Comparison of the Two Approaches:

Using the *Not In* clause with a subquery (Approach A) or using the *External Join* and checking if the *join* result is *null* (Approach B) didn't make any big difference in the execution times and costs.

Since both use a previously calculated *view*, either to check the *Not In* or to *Join*, they are going to accumulate that execution/calculation cost.

The major difference remains in the *where* conditions of the queries. In approach A the query only needs to check the entire *view* (previously calculated) for each record/row, but in approach B it only needs to check if the *join* to the other table was successful or not (check if the Id of the other table was filled). We think that the little execution time difference, between the two approaches, is caused by this.

3.4. Question 4

Description: Who was the professor with the most class hours for each class type in 2003/2004? Show the number and name of the professor, the kind of class, and the total number of class hours times the factor.

SQL Query:

```
SELECT MAX(Horas) AS Horas, MAX(NúmeroProfessor) AS NúmeroProfessor,
MAX(NomeProfessor) AS NomeProfessor, Tipo
FROM (
    SELECT SUM(dsd.horas) * dsd.fator AS Horas, tiposaula.tipo AS Tipo,
docentes.NR AS NúmeroProfessor, docentes.nome AS NomeProfessor
    FROM dsd
    LEFT JOIN tiposaula ON dsd.ID = tiposaula.ID
    LEFT JOIN docentes ON dsd.NR = docentes.NR
    WHERE tiposaula.ano_letivo = '2003/2004' AND dsd.horas IS NOT NULL AND
dsd.fator IS NOT NULL
    GROUP BY dsd.fator, tiposaula.tipo, docentes.NR, docentes.nome
    ORDER BY Horas DESC
)
GROUP BY Tipo
```

Retrieved Results:

	HORAS	NUMEROPROFESSOR	NOMEPROFESSOR	TIPO
1	30	908100	Álvaro Ferreira Marques Azevedo	P
2	3.5	246626	Raul Fernando de Almeida Moreira Vidal	OT
3	26	908290	Álvaro Ferreira Marques Azevedo	TP
4	30.67	909330	Álvaro Jorge da Maia Sêco	T

(figure 18 - Retrieved Results for Question 4)

Execution Data:

Environment X:

Query Execution Time: 0,035 seconds

OPERATION	OBJECT_NAME	OPTIONS	CARDINALITY	COST
SELECT STATEMENT				69
HASH		GROUP BY		5
VIEW				69
HASH		GROUP BY		5
HASH JOIN		RIGHT OUTER		69
Access Predicates XDSO.NR=XDOCENTES.NR(+)				68
TABLE ACCESS XDOCENTES	XDOCENTES	FULL		2567
HASH JOIN				5
Access Predicates XDSO.ID=XTIPOSAULA.ID				939
TABLE ACCESS XTIPOSAULA	XTIPOSAULA	FULL		2567
Filter Predicates XTIPOSAULA.ANO_LETIVO='2003/2004'				63
TABLE ACCESS XDSO	XDSO	FULL		1588
Filter Predicates XDSO.FATOR IS NOT NULL				36
				27

(figure 19 - Explain Plan of Environment X for Question 4)

Environment Y:

Query Execution Time: 0,032 seconds

OPERATION	OBJECT_NAME	OPTIONS	CARDINALITY	COST
SELECT STATEMENT				69
HASH		GROUP BY		5
VIEW				69
HASH		GROUP BY		5
HASH JOIN		RIGHT OUTER		69
Access Predicates YDSO.NR=YDOCENTES.NR(+)				68
TABLE ACCESS YDOCENTES	YDOCENTES	FULL		2567
HASH JOIN				5
Access Predicates YDSO.ID=YTIPOSAULA.ID				939
TABLE ACCESS YTIPOSAULA	YTIPOSAULA	FULL		2567
Filter Predicates YTIPOSAULA.ANO_LETIVO='2003/2004'				63
TABLE ACCESS YDSO	YDSO	FULL		1588
Filter Predicates YDSO.FATOR IS NOT NULL				36
				27

(figure 20 - Explain Plan of Environment Y for Question 4)

Environment Z:

Query Execution Time: 0,038 seconds

OPERATION	OBJECT_NAME	OPTIONS	CARDINALITY	COST
SELECT STATEMENT				69
HASH				5
VIEW		GROUP BY		69
HASH				2567
HASH JOIN		GROUP BY		69
Access Predicates		RIGHT OUTER		2567
ZDSD.NR=ZDOCENTES.NR(+)				68
TABLE ACCESS	ZDOCENTES	FULL		939
HASH JOIN				2567
Access Predicates				63
ZDSD.ID=ZTIPOSAULA.ID				
TABLE ACCESS	ZTIPOSAULA	FULL		1588
Filter Predicates				36
ZTIPOSAULA.ANO_LETIVO='2003/2004'				
TABLE ACCESS	ZDSD	FULL		27356
Filter Predicates				27
ZDSD.FATOR IS NOT NULL				

(figure 21 - Explain Plan of Environment Z for Question 4)

Comparison:

This comparison of execution environments was not conclusive at all. The cost of the query was the same in the three environments and even the operations executed for the query completion were the same, as we can see in the figures 19, 20 and 21.

The only difference we noticed was a little oscillation in the execution times, but it wasn't significant.

3.5. Question 5

Description: Compare the execution plans (only the environment Z) and the index sizes for the query, giving the course code, the academic year, the period, and the number of hours of the type 'OT' in the academic years of 2002/2003 and 2003/2004.

- With a B-tree index on the type and academic year columns of the ZTIPOSAULA table.
- With a bitmap index on the type and academic year columns of the ZTIPOSAULA table.

SQL Query:

```
SELECT ztiposaula.codigo, ztiposaula.ano_letivo, ztiposaula.periodo,
ztiposaula.horas_turno
FROM ztiposaula
WHERE ztiposaula.tipo = 'OT' AND (ztiposaula.ano_letivo = '2002/2003' OR
ztiposaula.ano_letivo = '2003/2004')
```

Retrieved Results:

	CODIGO	ANO_LETIVO	PERIODO	HORAS_TURNO
1	EIC5202	2002/2003	2S	0.5
2	EIC5202	2003/2004	2S	0.5

(figure 22 - Retrieved Results for Question 5)

Execution Data:

Index Approach A:

Query Execution Time: 0,002 seconds

OPERATION	OBJECT_NAME	OPTIONS	CARDINALITY	COST
SELECT STATEMENT				618
INLIST ITERATOR				30
TABLE ACCESS	ZTIPOSAULA	BY INDEX ROWID BATCHED		30
INDEX	BTREE_IDX_TIPO_AND_ANO_LETIVO	RANGE SCAN		5
Access Predicates				
AND				
ZTIPOSAULA.TIPO='OT'				
OR				
ZTIPOSAULA.ANO_LETIVO='2002/2003'				
ZTIPOSAULA.ANO_LETIVO='2003/2004'				

(figure 23 - Explain Plan of Approach A for Question 5)

Index Approach B:

Query Execution Time: 0,005 seconds

OPERATION	OBJECT_NAME	OPTIONS	CARDINALITY	COST
SELECT STATEMENT				618
TABLE ACCESS	ZTIPOSAULA	FULL		36
Filter Predicates				
AND				
ZTIPOSAULA.TIPO='OT'				
OR				
ZTIPOSAULA.ANO_LETIVO='2002/2003'				
ZTIPOSAULA.ANO_LETIVO='2003/2004'				

(figure 24 - Explain Plan of Approach B for Question 5)

Comparison:

A B-tree index is a balanced tree structure commonly used in databases to organize and search for keys in sorted order, making it efficient for range queries and disk-based systems. On the other hand, a Bitmap index represents a set of keys using a bitmap, providing space-efficient storage and fast set operations, particularly suitable for low cardinality data. The choice between them depends on factors such as the nature of the dataset and the types of queries to be performed.

Analyzing the data from the two executions (A (B-tree index) and B (Bitmap index)) we concluded that the B-tree index in approach A was a bit more efficient than the Bitmap index in approach B. The time in approach A was more than half of the time in B, and the cost for the B-tree index was 20% smaller than the cost for the Bitmap index.

With a cardinality greater than 600, and with the analysis we have done above, we think that a B-tree index is a better choice, in this case.

3.6. Question 6

Description: Select the programs (curso) with classes with all the existing types.

SQL Query:

```
SELECT Curso
FROM (
    SELECT COUNT(DISTINCT tiposaula.tipo) AS NumeroTipoAulas, ucs.curso AS
Curso
    FROM tiposaula
    LEFT JOIN ocorrencias ON tiposaula.codigo = ocorrencias.codigo AND
tiposaula.ano_letivo = ocorrencias.ano_letivo AND tiposaula.periodo =
ocorrencias.periodo
    LEFT JOIN ucs ON ocorrencias.codigo = ucs.codigo
    GROUP BY ucs.curso
)
WHERE NumeroTipoAulas = (
    SELECT COUNT (DISTINCT tiposaula.tipo)
    FROM tiposaula
)
```

Retrieved Results:

	CURSO
1	9461
2	4495
3	9508
4	2021

(figure 25 - Retrieved Results for Question 6)

Execution Data:

Environment X:

Query Execution Time: 0,046 seconds

OPERATION	OBJECT_NAME	OPTIONS	CARDINALITY	COST
SELECT STATEMENT				643
FILTER				2
Filter Predicates	COUNT(\$vm_col_1)=(SELECT COUNT(\$vm_col_1) FROM (SELECT XTIPOSAULA.TIPO \$vm_col_1 FROM XTIPOSAULA XTIPOSAULA GROUP BY XTIPOSAULA.TIPO) VM_NWWW_2)			
HASH		GROUP BY	2	643
VIEW	SYS.VM_NWWW_1		404	643
HASH		GROUP BY	404	643
HASH JOIN		RIGHT OUTER	21019	642
Access Predicates	XOCORRENCIAS.CODIGO=XUCS.CODIGO(+)			
TABLE ACCESS	XUCS	FULL	5396	13
HASH JOIN		RIGHT OUTER	21019	629
Access Predicates	XTIPOSAULA.CODIGO=XOCORRENCIAS.CODIGO(+) XTIPOSAULA.ANO_LETIVO=XOCORRENCIAS.ANO_LETIVO(+) XTIPOSAULA.PERIODO=XOCORRENCIAS.PERIODO(+)			
AND				
TABLE ACCESS	XOCORRENCIAS	FULL	21747	593
TABLE ACCESS	XTIPOSAULA	FULL	21019	36
SORT		AGGREGATE	1	37
VIEW	SYS.VM_NWWW_2		5	37
SORT		GROUP BY	5	37
TABLE ACCESS	XTIPOSAULA	FULL	21019	36
Other XML				

(figure 26 - Explain Plan of Environment X for Question 6)

Environment Y:

Query Execution Time: 0,055 seconds

OPERATION	OBJECT_NAME	OPTIONS	CARDINALITY	COST
SELECT STATEMENT				52
FILTER				2
Filter Predicates	COUNT(\$vm_col_1)=(SELECT COUNT(\$vm_col_1) FROM (SELECT YTIPOSAULA.TIPO \$vm_col_1 FROM YTIPOSAULA YTIPOSAULA GROUP BY YTIPOSAULA.TIPO) VM_NWWW_2)			
HASH		GROUP BY	2	52
VIEW	SYS.VM_NWWW_1		404	52
HASH		GROUP BY	404	52
HASH JOIN		RIGHT OUTER	21019	50
Access Predicates	YOCORRENCIAS.CODIGO=YUCS.CODIGO(+)			
TABLE ACCESS	YUCS	FULL	5396	13
NESTED LOOPS		OUTER	21019	37
TABLE ACCESS	YTIPOSAULA	FULL	21019	36
INDEX	YOCORRENCIAS_PK	UNIQUE SCAN	1	0
Access Predicates	YTIPOSAULA.CODIGO=YOCORRENCIAS.CODIGO(+) YTIPOSAULA.ANO_LETIVO=YOCORRENCIAS.ANO_LETIVO(+) YTIPOSAULA.PERIODO=YOCORRENCIAS.PERIODO(+)			
AND				
AGGREGATE			1	37
VIEW	SYS.VM_NWWW_2		5	37
SORT		GROUP BY	5	37
TABLE ACCESS	YTIPOSAULA	FULL	21019	36
Other XML				

(figure 27 - Explain Plan of Environment Y for Question 6)

Environment Z:

Query Execution Time: 0,051 seconds

OPERATION	OBJECT_NAME	OPTIONS	CARDINALITY	COST
SELECT STATEMENT				52
FILTER				2
Filter Predicates	COUNT(\$vm_col_1)=(SELECT COUNT(\$vm_col_1) FROM (SELECT ZTIPOSAULA.TIPO \$vm_col_1 FROM ZTIPOSAULA ZTIPOSAULA GROUP BY ZTIPOSAULA.TIPO) VM_NWWW_2)			
HASH		GROUP BY	2	52
VIEW	SYS.VM_NWWW_1			404
HASH		GROUP BY	404	52
HASH JOIN		RIGHT OUTER	21019	50
Access Predicates	ZOCORRENCIAS.CODIGO=ZUCS.CODIGO(+)			
TABLE ACCESS	ZUCS	FULL		5396
NESTED LOOPS		OUTER		21019
TABLE ACCESS	ZTIPOSAULA	FULL		21019
INDEX	ZOCORRENCIAS_PK	UNIQUE SCAN		1
Access Predicates	AND			
ZTIPOSAULA.CODIGO=ZOCORRENCIAS.CODIGO(+)				
ZTIPOSAULA.ANO_LETIVO=ZOCORRENCIAS.ANO_LETIVO(+)				
ZTIPOSAULA.PERIODO=ZOCORRENCIAS.PERIODO(+)				
SORT		AGGREGATE		1
VIEW	SYS.VM_NWWW_2			5
SORT		GROUP BY		5
INDEX	BTRREF_IDX_TIPO_AND_ANO_LETIVO_...	FAST FULL SCAN		21019

(figure 28 - Explain Plan of Environment Z for Question 6)

Comparison:

Analyzing the three environments explain plans, we can conclude that the addition of integrity constraints (primary keys (PK) and foreign keys (FK)) is essential for query performance, because the cost has reduced tremendously from the environment X to the environment Y.

We also noticed that the index applied in the table tiposaula has helped with the reduction of the cost too, although it was a very tiny improvement compared to the one made from the integrity constraints.

About the time execution, it was not very conclusive because the execution times collected were a lot similar, we only noticed a little oscillation in the execution times, but it wasn't significant.

4. Conclusions

Through comprehensive analysis of results, costs, and execution times across different query scenarios and execution environments, several key learning lessons and takeaways have emerged.

The study on query optimization revealed the importance of effective query organization, indexing, and query restructuring in improving database performance. By aligning query structures with the data model and using appropriate indexing techniques, execution times were significantly improved. Environmental factors, such as system resources and concurrent workloads, also influence execution statistics. Indexing plays a crucial role in query performance, reducing execution times but requiring a balance between utilization and maintenance costs.